



# NIST AI RMF Implementation Playbook

**Project:** Secure AI Gateway Implementation

**Framework:** NIST AI Risk Management Framework (Functions: Manage, Measure)

**Infrastructure:** Cloudflare AI Gateway, Workers AI (Llama-3), Python

---

## Phase 1: The Objective

To move from "Direct API Access" (High Risk, No Visibility) to a **Managed AI Pipeline** (Low Risk, High Visibility).

- **Goal:** Intercept, log, and control all traffic sent to Large Language Models (LLMs).
  - **NIST Alignment:**
    - **Manage:** Control access via Gateway authentication.
    - **Measure:** Log prompts, tokens, and latency for auditability.
- 

## Phase 2: Account Discovery (The "ID Trap")

Before writing code, we had to locate the correct infrastructure identifiers. *Crucial Lesson: Do not confuse Zone IDs with Account IDs.*

### Step 1: Locate Account ID

1. Log in to [Cloudflare Dashboard](#).
  2. Look at the browser URL: [https://dash.cloudflare.com/YOUR\\_ACCOUNT\\_ID/ai/gateway](https://dash.cloudflare.com/YOUR_ACCOUNT_ID/ai/gateway)
  3. **Artifact:** The 32-character string after .com/ is the **Account ID**.
    - **Correct:** c267... (From URL)
    - **Incorrect:** Any ID found on the "Website Overview" sidebar (Zone ID).
-

## Phase 3: Authentication Strategy

We attempted "Least Privilege" (Tokens) but fell back to "Administrator Access" (Global Key) due to scope friction.

### Method A: API Tokens (Preferred for Prod)

- **Path:** User Profile > API Tokens > Create Token.
- **Permissions Required:**
  - Account > AI Gateway > Run
  - Account > Workers AI > Read
- **The Critical Fix:** You must select "**All Accounts**" under *Account Resources* to prevent scoping errors (Error 10000).

### Method B: Global API Key (The "Master Key")

- **Used For:** Debugging when Tokens fail with Authentication Error (10000).
- **Path:** User Profile > API Tokens > API Keys > Global API Key.
- **Headers Used:** X-Auth-Email and X-Auth-Key.

---

## Phase 4: Gateway Infrastructure Setup

We established the control plane.

### Step 1: Create the Gateway

1. Navigate to AI > **AI Gateway**.
2. Click **Create Gateway**.
3. **Name:** secure-scanner-gateway.
4. **Action:** This generates a "Universal Endpoint" URL:  
`https://gateway.ai.cloudflare.com/v1/{ACCOUNT_ID}/{GATEWAY_NAME}`

### Step 2: The "Unlock" (Crucial Fix)

- **Issue:** The Gateway was created with "Authentication Enabled" (Lock Icon ).
- **Symptom:** All API calls returned 403 Forbidden or Unauthorized.
- **The Fix:**
  1. Click secure-scanner-gateway.
  2. Go to **Settings**.
  3. **Toggle OFF** "API Authentication" (Enforce Access Control).
  4. **Verify:** The Lock Icon  must disappear from the dashboard list.

# Phase 5: The Implementation (Python)

We built a script to route prompts through the Gateway instead of hitting the model directly.

## The Final Artifact (ai\_ops\_test.py)

This is the working code block used to validate the pipeline.

Python

```
import requests

# --- CONFIGURATION ---
GATEWAY_NAME = "secure-scanner-gateway"
ACCOUNT_ID = "c267966add7c04e16295a38bef315e66"
EMAIL = "tharunrs007@gmail.com"
GLOBAL_KEY = "PASTE_YOUR_GLOBAL_KEY_HERE" # <--- The Master Key
# -----


def run_pipeline():
    print(f"🚀 Targeting Gateway: {GATEWAY_NAME}")

    # The Gateway URL (Not the Dashboard URL!)
    url =
        f"https://gateway.ai.cloudflare.com/v1/{ACCOUNT_ID}/{GATEWAY_NAME}/workers-ai/@cf/meta/llama-3-8b-instruct"

    headers = {
        "X-Auth-Email": EMAIL,
        "X-Auth-Key": GLOBAL_KEY.strip(),
        "Content-Type": "application/json"
    }

    payload = {
        "messages": [
            {"role": "system", "content": "You are a specific security bot."},
            {"role": "user", "content": "Confirm system operational."}
        ]
    }

    try:
```

```

response = requests.post(url, headers=headers, json=payload)
if response.status_code == 200:
    print("✅ VICTORY! Response received through Gateway.")
    print(f"🤖 AI: {response.json()['result']['response']}")
else:
    print(f"❌ Error {response.status_code}: {response.text}")
except Exception as e:
    print(f"❌ Script Failed: {e}")

if __name__ == "__main__":
    run_pipeline()

```

---

## Phase 6: Verification (NIST "Measure")

We confirmed the system works by checking the telemetry.

1. **Action:** Run the Python script.
2. **Verify:** Go to Cloudflare Dashboard > AI Gateway > secure-scanner-gateway > **Logs**.
3. **Evidence:**
  - o **Prompt:** Visible ("Confirm system operational").
  - o **Status:** Success (200 OK).
  - o **Metadata:** Token count and Latency are recorded.

---

## Appendix: Troubleshooting Log

*This section records the errors we faced so we recognize them next time.*

Error Code	Meaning	Root Cause	The Fix
<b>10000</b>	Authentication Error	Token Scope Mismatch	Used "Global API Key" or recreated token with "All Accounts" scope.
<b>401</b>	Unauthorized	Wrong Credentials	Verified Account ID vs. Zone ID; switched to Global Key.

<b>2001</b>	Gateway Not Configured	Script tried to use a non-existent gateway	Created Gateway manually in Dashboard to accept TOS.
<b>403</b>	Forbidden	Gateway Locked 	Turned OFF "API Authentication" in Gateway Settings.
<b>400</b>	Bad Request	URL Error	Switched from dash.cloudflare.com URL to gateway.ai.cloudflare.com.

---

## Phase 7: Availability Engineering (DoS Defense)

- **Objective:** Prevent resource exhaustion attacks by limiting request volume.
- **Configuration:**
  - **Mechanism:** Cloudflare Rate Limiting (Fixed Window).
  - **Policy:** 3 requests / 1 minute (Test Mode).
  - **Outcome:** Requests >3 receive **HTTP 429**.
- **Validation:**
  - Developed `dos_attack_test.py` to simulate high-frequency traffic.
  - Validated that the Gateway fails open (200 OK) for legitimate traffic and fails closed (429 Blocked) for abuse.

## Phase 8: Confidentiality Engineering (Data Privacy)

- **Objective:** Prevent data leakage (PII, Credentials, Secrets) into persistent storage logs, ensuring compliance with privacy standards (GDPR/CCPA).
  - **Configuration:**
    - **Mechanism:** Cloudflare Log Redaction.
    - **Action:** Disabled "*Store request and response payloads*" in Gateway Settings.
    - **Outcome:** Gateway processes prompts for metrics (tokens/cost) but **discards** the actual text content immediately.
  - **Validation:**
    - Developed `privacy_test.py` to inject a "Canary Token" (fake secret password).
    - **Verified:** The Gateway processed the request (200 OK), but the Logs showed [Redacted] or empty fields for the message content.
- 

## Phase 9: Final Architecture Summary

- **Current State:**
  - **Access:** Locked via API Tokens/Global Keys (**Integrity**).
  - **Traffic:** Routed through `secure-scanner-gateway` (**NIST Manage**).
  - **Resilience:** Rate Limited to 3 req/min to block DoS attacks (**Availability**).
  - **Privacy:** Payload logging disabled to protect secrets (**Confidentiality**).
  - **Observability:** Full visibility into Token Usage, Latency, and Error Rates (**NIST Measure**).